

1. Data split & validation

Train: Jan-Oct 2025 (48,000 rows). Final test: Nov-Dec 2025 (12,000 rows, labels hidden). Because market_index drifts (train mean 1.08 vs validation 0.93) and rates trend by month, a random split would leak the future and give optimistic scores. We use a temporal expanding window with 3 folds:
Fold1: train Jan-Apr -> test May-Jun
Fold2: train Jan-Jun -> test Jul-Aug
Fold3: train Jan-Aug -> test Sep-Oct (closest proxy for Nov-Dec).
Metrics per fold: MAE / RMSE / MAPE / MedAPE. Model selection on mean MAE/MAPE.

- weight: 300 NaN (flag + median), 292 negative -> abs() (sign-entry error), 1204 capped at 47500 (flag).
- market_index: 374 NaN train / 249 val -> same-day median, fallback global median (+ missing flag). Within-day noise ~0.025 so daily median denoises safely.
- distance: 48 rows floored at 70.0 (flag).
- posted_rate: ~1.4% outliers (x3 or /3) -> kept; robust L1 loss on log-target handles them.
- coordinates are not real geography (e.g. Phoenix shifted); used as lane embeddings, not maps.
- 8 unseen validation cities (Allentown, Charlotte, Chicago, Jackson, Knoxville, Laredo, Norfolk, San Diego): no city-ID features, coordinates only.

- distance <-> rate corr 0.91: the dominant driver.
- market_index has a huge weekly cycle (Thu ~1.18, Sun ~0.99) but moves prices only ~+-1%%; raw index is a trap, daily-median smoothing (mi_smooth) is used.
- quote_signal is U-shaped vs \$/mile (edges 3.2-3.4, middle 2.1): strong non-linear signal, qs_dev=|q-median| added.
- monthly \$/mile drifts 2.10 (Jan) -> 2.33 (Jun) -> 2.24 (Oct): independent time trend + trend/dow features.

LightGBM regression_l1 on log1p(posted_rate). GBM captures the U-shaped quote effect and lane interactions that linear models miss; log-target stabilizes skew; L1 ignores label outliers (RMSE ~630 for all models is outlier-driven). Coordinates + 3-way equipment one-hot generalize to new cities.

fold	model	MAE	RMSE	MAPE	MedAPE
1	naive	206.18	654.18	11.01	7.05
1	ridge-log	104.33	631.53	4.31	1.75
1	hgb-logL2	135.42	641.05	5.79	2.85
1	lgbm-logL1	119.43	636.45	4.83	1.82
2	naive	222.73	662.87	10.11	6.70
2	ridge-log	141.03	624.39	6.01	3.42
2	hgb-logL2	148.45	628.69	6.33	3.28
2	lgbm-logL1	121.96	624.70	5.31	2.20
3	naive	229.10	670.66	10.49	6.80
3	ridge-log	110.64	634.70	4.78	1.97
3	hgb-logL2	119.12	637.75	5.11	2.25
3	lgbm-logL1	106.48	633.92	4.66	1.90

Mean over folds:

	MAE	RMSE	MAPE	MedAPE
model				
hgb-logL2	134.33	635.83	5.74	2.79
lgbm-logL1	115.96	631.69	4.94	1.97
naive	219.34	662.57	10.53	6.85
ridge-log	118.67	630.20	5.03	2.38

Winner: lgbm-logL1 (mean MAE 115.96, MAPE 4.94) vs ridge-log (118.67 / 5.03), hgb-logL2 (134.33 / 5.74), naive median \$/mile (219.34 / 10.53). Fold3 (Sep-Oct, best Nov-Dec proxy): lgbm MAE 106.5 / MAPE 4.66. Final model

Freight Rate Prediction - Spotter ML Assessment

retrained on all Jan-Oct with a time-tail eval split + early stopping (seed 0).


6. De

Lane Lexington -> Fort Wayne, 360 mi, Dry Van, 32,000 lb; only date varies. december_chart_inputs.csv ships WITHOUT market_index / quote_signal / coordinates, so:

- coords: train city medians (both cities exist in train).
- market_index + mi_smooth: calendar-day means from validation.csv (input features covering all Dec dates, not labels).
- quote_signal: December median from validation (2.0512).

Prediction range ~\$819-826 with a small weekly cycle on top of the trend level. Limitation: 10 months cannot separate trend from seasonality, and holiday December is out-of-distribution; chart shape is more reliable than absolute level.

Candidate: December 2025 Predicted Load Rate



7. Re

```
pip install -r requirements.txt
python -m src.validate
python -m src.predict
python score.py --predictions validation_predictions.csv --december-predictions data/december_chart_inputs.csv
```